

# Notes for Time-Series and Machine Learning Methods

## Part A: Exponential Smoothing and SARIMAX

Natarajan Gautam

EECS, Syracuse University

# Time-Series Data

- ▶ We collect data from observations over time
- ▶ We denote the observation numbers (i.e., *index*) as  $1, 2, \dots, T$
- ▶ What we observe is denoted as  $y_1, y_2, \dots, y_T$  (see below)
- ▶ So  $y_t$  for  $t \in \{1, \dots, T\}$  is the  $t^{\text{th}}$  observation
- ▶ For example, the monthly average number of calls, daily average sales per store across locations, amount of solar power generated, etc. (see below)
- ▶ We wish to forecast  $y_{T+h}$  for some pre-defined horizon  $h$  which is a positive integer
- ▶ That  $h$ -step ahead forecast for observation  $T + h$  given we have observed values from time 1 to  $T$  is represented as  $\hat{y}_{T+h|T}$
- ▶ Below is a 15-minute average solar irradiance dataset sample  $y$

Index	Date	Time	Avg. Irradiance
1	Apr-5-2023	11:00-11:14	632 = $y_1$
2	Apr-5-2023	11:15-11:29	641 = $y_2$
3	Apr-5-2023	11:30-11:44	579 = $y_3$
$\vdots$	$\vdots$	$\vdots$	$\vdots$
$T$	Apr-12-2023	11:15-11:29	645 = $y_T$

# Time-Series Forecast: Exponential Smoothing (ES)

- ▶ Why at all exponential smoothing?
  - ▶ Handles seasonal data well
  - ▶ Unlike SARIMAX, ES very rarely throws errors
  - ▶ A special case is indeed persistent forecast
- ▶ We begin with the simple time series without any trend or seasonality
- ▶ A one-step ahead forecast for time  $T + 1$  given we have observed values from time 1 to  $T$  is represented as  $\hat{y}_{T+1|T}$
- ▶ The most fundamental exponential smoothing method uses

$$\hat{y}_{T+1|T} = \alpha y_T + \alpha(1 - \alpha)y_{T-1} + \alpha(1 - \alpha)^2 y_{T-2} + \dots$$

where  $\alpha \in [0, 1]$  is called the smoothing parameter

- ▶ Note that larger weights are given to most recent observations
- ▶ The weights attached to the observations decrease exponentially as we go back in time, hence the name ES

$\hat{y}_{T+1|T} = y_T$   
is persistent forecast

$h=1$

$\alpha = 1 \Rightarrow$  persistent forecast

We do not know  $\alpha$   
 $y_0$  is unknown

## ES: Component Form for Simple Time Series

$$\hat{y}_{T+1|T} = \alpha y_T + \alpha(1-\alpha)y_{T-1} + \alpha(1-\alpha)^2 y_{T-2} + \dots$$

For any  $t$

$$\hat{y}_{t+1|t} = \alpha y_t + \alpha(1-\alpha)y_{t-1} + \alpha(1-\alpha)^2 y_{t-2} + \dots$$

- ▶ We can show that for any  $t$ ,

$$\hat{y}_{t+1|t} = \alpha y_t + (1-\alpha)\hat{y}_{t|t-1}$$

$$\hat{y}_{t+1|t} = \alpha y_t + (1-\alpha) \left[ \alpha y_{t-1} + \alpha(1-\alpha)y_{t-2} + \dots \right]$$

$\uparrow$   
 $\hat{y}_{t|t-1}$

- ▶ The above is written in component form as
  - ▶ Forecast equation:  $\hat{y}_{t+1|t} = \ell_t$
  - ▶ Smoothing equation:  $\ell_t = \alpha y_t + (1-\alpha)\ell_{t-1}$
- ▶ We call  $y_1, y_2, \dots, y_T$  as the training data
- ▶ There are two unknowns  $\alpha$  and  $\ell_0$
- ▶ We select  $\alpha$  and  $\ell_0$  by minimizing the sum of squared errors (SSE)

$$SSE = \sum_{t=1}^T (y_t - \hat{y}_{t|t-1})^2 //$$

- ▶ The minimization is usually done with an optimization tool

## ES: Component Form for Time Series with Trend but no Seasonality

### ▶ Holt's linear trend method

▶ Forecast equation:  $\hat{y}_{t+1|t} = \ell_t + b_t$

▶ Level equation:  $\ell_t = \alpha y_t + (1 - \alpha)(\ell_{t-1} + b_{t-1})$

▶ Trend equation:  $b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}$

$\ell_t$  is an estimate of the level of the series at time  $t$ ,  $b_t$  is an estimate of the trend of the series at time  $t$ ,  $\alpha$  the smoothing parameter for level, and  $\beta^*$  the smoothing parameter for trend

▶ We could consider other aspects such as damped trend, and multiplicative trend (above is additive)

▶ See Hyndman and Athanasopoulos' FPP3 book for details (same notation)

⇒ ▶ We will not use damping in this course but is an effective tool that prevents extremely high forecast values

▶ In fact, this case of no seasonality itself is something we do not consider in this course

## ES: Component Form for Time Series with Trend and Seasonality

- ▶ Results for any horizon  $h$  (earlier we had  $h = 1$ )
- ▶ The seasonality period is  $m$  and define  $k = \lfloor (h - 1)/m \rfloor = 0$
- ▶ Holt-Winters' seasonal method uses the iteration
  - ▶ Forecast equation:  $\hat{y}_{t+h|t} = \ell_t + hb_t + s_{t+h-m(k+1)}$
  - ▶ Level equation:  $\ell_t = \alpha(y_t - s_{t-m}) + (1 - \alpha)(\ell_{t-1} + b_{t-1})$
  - ▶ Trend equation:  $b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}$
  - ▶ Seasonal equation:  $s_t = \gamma(y_t - \ell_{t-1} - b_{t-1}) + (1 - \gamma)s_{t-m}$


*if  $m > h-1$   
i.e. predict less than a seasonal period*

- ▶ Multiplicative seasonality (useful for solar forecast)

- ▶ Forecast equation:  $\hat{y}_{t+h|t} = (\ell_t + hb_t)s_{t+h-m(k+1)}$
- ▶ Level equation:  $\ell_t = \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha)(\ell_{t-1} + b_{t-1})$
- ▶ Trend equation:  $b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}$
- ▶ Seasonal equation:  $s_t = \gamma \frac{y_t}{\ell_{t-1} + b_{t-1}} + (1 - \gamma)s_{t-m}$

- ▶ We select  $\alpha, \beta^*, \gamma, s_0, s_{-1} \dots, s_{-m+1}, b_0$  and  $\ell_0$  by minimizing the RMSE or maximizing the likelihood *→ Python code*
- ▶ Software also allows for a Box-Cox transformation *→ Make data normal*

## ES: Python Code



- ▶ We use `statsmodels.tsa.api` for ES by importing `ExponentialSmoothing`
- ▶ Refer to the documentation to see what all parameters can be set
- ▶ Check what the default values are for them
- ▶ In our case we may have historical data that is much larger than  $T$
- ✓ ▶ We select  $T$ , train the model, and predict for  $T + h$
- ✓ ▶ Next time we use data from index 2 to  $T + 1$ , again train the model, and predict for  $T + 1 + h$
- ✓ ▶ We continue this process of sliding window
- ⇒ ▶ We let the model choose  $\alpha$ ,  $\beta^*$ ,  $\gamma$ ,  $s_0$ ,  $s_{-1} \dots, s_{-m+1}$ ,  $b_0$  and  $\ell_0$  by the default optimization method
- ▶ Technically speaking we may already know some of them as we are considering a sliding window
- ▶ We ignore that aspect but note that these can be provided to the function
- ▶ For point forecasts we will not do a Box-Cox transformation
- ▶ Let us use the code provided in the course







## Time-Series Forecasts: SARIMAX

- ▶ We begin by explaining ARIMA (Auto-Regressive Integrated Moving Average), i.e., without seasonality and exogenous data
- ▶ The approach taken in ARIMA is to first remove the effects of seasonality and trend using a process called differencing
- ▶ Then ARIMA simply uses recency to make the forecasts, and then re-introduce the seasonality and trend
- ▶ Hence, as an initial step, we try to create a *stationary* time-series
- ▶ A stationary time-series in an ideal condition is independent of the time stamp (index)
- ▶ Both trend and seasonality would render a time-series to be non-stationary and they would have to be removed
- ▶ ARIMA assumes the time-series to be stationary
- ▶ The approach used to create a “nearly” stationary time-series is called differencing

# Differencing

- ▶ Recall that we have observations  $y_1, y_2, \dots, y_T$  that we use for making the forecasts
- ▶ We denote the time series as  $\{y_t : 1 \leq t \leq T\}$
- ▶ Now consider the time-series  $\{y'_t : 2 \leq t \leq T\}$  which is called the first-order differencing term defined as

$$y'_t = y_t - y_{t-1}$$

for all  $t \in \{2, 3, \dots, T\}$

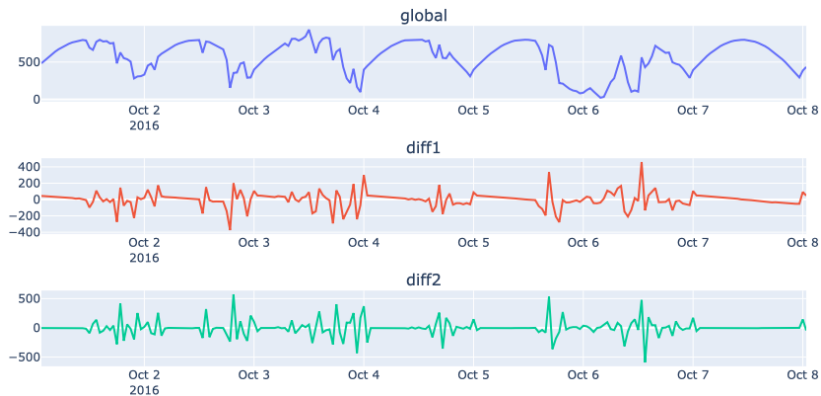
- ▶ The differenced series has one fewer term than the original time-series
- ▶ Sometimes the series  $\{y'_t : 1 \leq t \leq T\}$  may not be adequately stationary, and it may be required to do one more round of differencing
- ▶ For that, consider the time series  $\{y''_t : 2 \leq t \leq T\}$  defined as

$$y''_t = y'_t - y'_{t-1}$$

for all  $t \in \{3, 4, \dots, T\}$

- ▶ Note that the second order differenced series has one fewer term than the first-order time-series and two fewer than the original time series

# Differencing Example



## Differencing: Some Comments and Seasonal Differencing

- ▶ In the previous figure, the automated code using `pmdarima` picks first order
- ▶ Usually first order differencing is enough, rarely second order
- ▶ It is not recommended to go beyond second order differencing, although the series may not be stationary
- ▶ Sometimes that is because of lingering seasonal effects
- ▶ When there are seasonal effects, it may be worthwhile (although not guaranteed to remove) to perform seasonal differencing
- ▶ Let  $\{y_t^s : m \leq t \leq T\}$  be lag- $m$  differences, where  $m$  is the period of seasonality
- ▶ The seasonal differencing term is defined as

$$y_t^s = y_t - y_{t-m}$$

for all  $t \in \{m+1, m+2, \dots, T\}$

- ▶ If needed, we could perform a second-differences (one against season and one against consecutive terms)

## Time-Series Forecasts: Model for ARIMA

- ▶ We begin with ARIMA with no seasonal effects.
- ▶ We stationarize the time series using differencing
- ▶ The **ARIMA**( $p, d, q$ ) model has  $p$  as the autoregressive part,  $d$  the degree of differencing and  $q$  the moving average part
- ▶ This gives rise to the model (in terms of the differenced series  $y'_t$ )

$$y'_t = c + \phi_1 y'_{t-1} + \dots + \phi_p y'_{t-p} + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q} + \epsilon_t$$

where  $\epsilon_t$  is from the original series

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \epsilon_t$$

- ▶ Also, the  $\epsilon_t$  terms are a function of  $c, \phi_1, \dots, \phi_p$  as well as the time series
- ▶ Note that  $c, \phi_1, \dots, \phi_p, \theta_1, \dots, \theta_q$  are all unknown

## Time-Series Forecasts: Estimating the ARIMA Unknowns

- ▶ In the ARIMA model we can use MLE (default in the python library we use) to estimate  $c, \phi_1, \dots, \phi_p, \theta_1, \dots, \theta_q$  which are all unknown
- ▶ For the MLE, we maximize the likelihood of obtaining the series  $y'_t$  for all  $t$  assuming it is Gaussian and so are the  $\epsilon_t$  values
- ▶ Another approach is to minimize the sum of squares of  $\epsilon_t$
- ▶ Either case would result in a complicated non-linear optimization problem that cannot be written in closed-form
- ▶ There are numerous software packages that produce the results

## Time-Series Forecasts: Model for SARIMAX

- ▶ The key idea of seasonal-ARIMA is to add seasonal terms
- ▶ Extend the **ARIMA**( $p, d, q$ ) model to the **seasonal-ARIMA**( $p, d, q$ )( $P, D, Q$ ) with  $m$  seasons
- ▶ Although we do not present the equation, it is worthwhile understanding that we add seasonality AR terms  $y_{t-m}, \dots, y_{t-Pm}$  with coefficients  $\Phi_1, \dots, \Phi_P$ , as well as seasonality MA terms  $\epsilon_{t-m}, \dots, \epsilon_{t-Qm}$  with coefficients  $\Theta_1, \dots, \Theta_Q$
- ▶ Now, if we were to add exogenous variables, say  $x_1, x_2, \dots, x_r$ , then we essentially get the form

$$y'_t = c + b_1x_1 + b_2x_2 + \dots + b_rx_r + \phi_1y'_{t-1} + \dots$$

## Selecting Model Parameters for SARIMAX

- ▶ Next we need to determine how to get  $p$ ,  $q$ ,  $P$ , and  $Q$
- ▶ They can be selected by searching through the space of values and find the one that minimizes an information criterion
- ▶ The common choice of information criteria are AIC (Akaike's Information Criterion), AICc (i.e. AIC corrected), and BIC (Bayesian Information Criterion), see Hyndman and Athanasopoulos' FPP3 book
- ▶ We do not present the expressions here but they are typically functions of likelihood of obtain the data (that was used in MLE),  $T$ ,  $p$ ,  $q$ ,  $P$ ,  $Q$ , and whether or not  $c$  was used
- ▶ There is also an automated way (although many do not recommend) to do this using `pmdarima`



## SARIMAX Forecasts: Python Code

- ▶ Import various python tools and libraries
- ▶ Select params viz. test data, sample size, etc.
- ▶ Get data and aggregate, if necessary
- ▶ Clean up data by imputing
- ▶ Merge with exogenous data (if we use them) to get required dataset
- ▶ Visualize the data for a few days
- ▶ Fit auto-arima model, or other methods for  $(p, d, q)$  and  $(P, D, Q)$
- ▶ Use sliding window for predictions
- ▶ Test predictions using various metrics