

Notes for Time-Series and Machine Learning Methods

Part A: Exponential Smoothing and SARIMAX

Last class

↓
Today

Natarajan Gautam

EECS, Syracuse University

Time-Series Data

- ▶ We collect data from observations over time
- ▶ We denote the observation numbers (i.e., *index*) as $1, 2, \dots, T$
- ▶ What we observe is denoted as y_1, y_2, \dots, y_T
- ▶ So y_t for $t \in \{1, \dots, T\}$ is the t^{th} observation
- ▶ For example, the monthly average number of calls, daily average sales per store across locations, amount of solar power generated, etc.
- ▶ We wish to forecast y_{T+h} for some pre-defined horizon h which is a positive integer
- ▶ That h -step ahead forecast for observation $T + h$ given we have observed values from time 1 to T is represented as $\hat{y}_{T+h|T}$
- ▶ Below is a 15-minute average solar irradiance dataset sample

Index	Date	Time	Avg. Irradiance
1	Apr-5-2023	11:00-11:14	632
2	Apr-5-2023	11:15-11:29	641
3	Apr-5-2023	11:30-11:44	579
\vdots	\vdots	\vdots	\vdots
T	Apr-12-2023	11:15-11:29	645

forecast

y_1
 y_2
 y_3
 \vdots
 y_T

$y_{T+h} = ?$

\hat{y}_{T+h}

Time-Series Forecast: Exponential Smoothing (ES)

- ▶ Why at all exponential smoothing?
 - ▶ Handles seasonal data well
 - ▶ Unlike SARIMAX, ES very rarely throws errors
 - ▶ A special case is indeed persistent forecast
- ▶ We begin with the simple time series **without** any trend or seasonality
- ▶ A one-step ahead forecast for time $T + 1$ given we have observed values from time 1 to T is represented as $\hat{y}_{T+1|T}$
- ▶ The most fundamental exponential smoothing method uses

$$\hat{y}_{T+1|T} = \alpha y_T + \alpha(1 - \alpha)y_{T-1} + \alpha(1 - \alpha)^2 y_{T-2} + \dots$$

where $\alpha \in [0, 1]$ is called the smoothing parameter

- ▶ Note that larger weights are given to most recent observations
- ▶ The weights attached to the observations decrease exponentially as we go back in time, hence the name ES

ES: Component Form for Simple Time Series

- ▶ We can show that for any t ,

$$\hat{y}_{t+1|t} = \alpha y_t + (1 - \alpha)\hat{y}_{t|t-1}$$

- ▶ The above is written in component form as

- ▶ Forecast equation: $\hat{y}_{t+1|t} = \ell_t$

- ▶ Smoothing equation: $\ell_t = \alpha y_t + (1 - \alpha)\ell_{t-1}$

- ▶ We call y_1, y_2, \dots, y_T as the training data

- ▶ There are two unknowns α and ℓ_0

- ▶ We select α and ℓ_0 by minimizing the sum of squared errors (SSE)

$$SSE = \sum_{t=1}^T (y_t - \hat{y}_{t|t-1})^2$$

- ▶ The minimization is usually done with an optimization tool

ES: Component Form for Time Series with Trend but no Seasonality

▶ Holt's linear trend method

- ▶ Forecast equation: $\hat{y}_{t+1|t} = \ell_t + b_t$
- ▶ Level equation: $\ell_t = \alpha y_t + (1 - \alpha)(\ell_{t-1} + b_{t-1})$
- ▶ Trend equation: $b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}$

ℓ_t is an estimate of the level of the series at time t , b_t is an estimate of the trend of the series at time t , α the smoothing parameter for level, and β^* the smoothing parameter for trend

- ▶ We could consider other aspects such as *damped* trend, and multiplicative trend (above is additive)
- ▶ See Hyndman and Athanasopoulos' FPP3 book for details
- ▶ We will not use damping in this course but is an effective tool that prevents extremely high forecast values
- ▶ In fact, this case of no seasonality itself is something we do not consider in this course

ES: Component Form for Time Series with Trend and Seasonality

- ▶ Results for any horizon h (earlier we had $h = 1$)
- ▶ The seasonality period is m and define $k = \lfloor (h - 1)/m \rfloor$
- ▶ Holt-Winters' seasonal method uses the iteration
 - ▶ Forecast equation: $\hat{y}_{t+h|t} = \ell_t + hb_t + s_{t+h-m(k+1)}$
 - ▶ Level equation: $\ell_t = \alpha(y_t - s_{t-m}) + (1 - \alpha)(\ell_{t-1} + b_{t-1})$
 - ▶ Trend equation: $b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}$
 - ▶ Seasonal equation: $s_t = \gamma(y_t - \ell_{t-1} - b_{t-1}) + (1 - \gamma)s_{t-m}$
- ▶ Multiplicative seasonality (useful for solar forecast)
 - ▶ Forecast equation: $\hat{y}_{t+h|t} = (\ell_t + hb_t)s_{t+h-m(k+1)}$
 - ▶ Level equation: $\ell_t = \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha)(\ell_{t-1} + b_{t-1})$
 - ▶ Trend equation: $b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}$
 - ▶ Seasonal equation: $s_t = \gamma \frac{y_t}{\ell_{t-1} + b_{t-1}} + (1 - \gamma)s_{t-m}$
- ▶ We select $\alpha, \beta^*, \gamma, s_0, s_{-1} \dots, s_{-m+1}, b_0$ and ℓ_0 by minimizing the RMSE or maximizing the likelihood
- ▶ Software also allows for a Box-Cox transformation

ES: Python Code

- ▶ We use `statsmodels.tsa.api` for ES by importing `ExponentialSmoothing`
- ▶ Refer to the documentation to see what all parameters can be set
- ▶ Check what the default values are for them
- ▶ In our case we may have historical data that is much larger than T
- ▶ We select T , train the model, and predict for $T + h$
- ▶ Next time we use data from index 2 to $T + 1$, again train the model, and predict for $T + 1 + h$
- ▶ We continue this process of sliding window
- ▶ We let the model choose α , β^* , γ , s_0 , $s_{-1} \dots, s_{-m+1}$, b_0 and ℓ_0 by the default optimization method
- ▶ Technically speaking we may already know some of them as we are considering a sliding window
- ▶ We ignore that aspect but note that these can be provided to the function
- ▶ For point forecasts we will not do a Box-Cox transformation
- ▶ Let us use the code provided in the course

2/25
Quiz 3
Forecasting


2/27
Practice
midterm

3/4
In-class
midterm


Time-Series Forecasts: SARIMAX


- ▶ We begin by explaining ARIMA (Auto-Regressive Integrated Moving Average), i.e., without seasonality and exogenous data
- ▶ The approach taken in ARIMA is to first remove the effects of seasonality and trend using a process called differencing
- ▶ Then ARIMA simply uses recency to make the forecasts, and then re-introduce the seasonality and trend
- ▶ Hence, as an initial step, we try to create a stationary time-series
- ▶ A stationary time-series in an ideal condition is independent of the time stamp (index)
- ▶ Both trend and seasonality would render a time-series to be non-stationary and they would have to be removed
- ▶ ARIMA assumes the time-series to be stationary
- ▶ The approach used to create a “nearly” stationary time-series is called differencing

Differencing

- ▶ Recall that we have observations y_1, y_2, \dots, y_T that we use for making the forecasts
- ▶ We denote the time series as $\{y_t : 1 \leq t \leq T\}$ 
- ▶ Now consider the time-series $\{y'_t : 2 \leq t \leq T\}$ which is called the first-order differencing term defined as


$$y'_t = y_t - y_{t-1}$$

 1st differencing

for all $t \in \{2, 3, \dots, T\}$ 

- ▶ The differenced series has one fewer term than the original time-series
- ▶ Sometimes the series $\{y'_t : 1 \leq t \leq T\}$ may not be adequately stationary, and it may be required to do one more round of differencing
- ▶ For that, consider the time series $\{y''_t : 2 \leq t \leq T\}$ defined as

$$y''_t = y'_t - y'_{t-1}$$

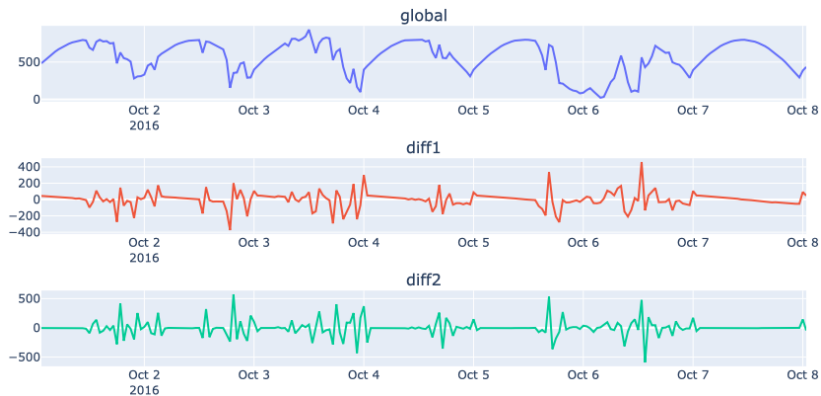
 2nd differencing

for all $t \in \{3, 4, \dots, T\}$

- ▶ Note that the second order differenced series has one fewer term than the first-order time series and two fewer than the original time series

 differencing

Differencing Example



Differencing: Some Comments and Seasonal Differencing

- ▶ In the previous figure, the automated code using `pmdarima` picks first order
- ▶ Usually first order differencing is enough, rarely second order
- ▶ It is not recommended to go beyond second order differencing, although the series may not be stationary
- ▶ Sometimes that is because of lingering seasonal effects
- ▶ When there are seasonal effects, it may be worthwhile (although not guaranteed to remove) to perform seasonal differencing
- ▶ Let $\{y_t^s : m \leq t \leq T\}$ be lag- m differences, where m is the period of seasonality
- ▶ The seasonal differencing term is defined as

$$y_t^s = y_t - y_{t-m}$$

for all $t \in \{m+1, m+2, \dots, T\}$

- ▶ If needed, we could perform a second-differences (one against season and one against consecutive terms)

Time-Series Forecasts: Model for ARIMA

- ▶ We begin with ARIMA with no seasonal effects.
- ▶ We stationarize the time series using differencing
- ▶ The **ARIMA**(p, d, q) model has p as the autoregressive part, d the degree of differencing and q the moving average part
- ▶ This gives rise to the model (in terms of the differenced series y'_t)

$d = 0 \text{ or } 1 \text{ or } 2$

$$y'_t = c + \phi_1 y'_{t-1} + \dots + \phi_p y'_{t-p} + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q} + \epsilon_t$$

where ϵ_t is from the original series

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \epsilon_t$$


- ▶ Also, the ϵ_t terms are a function of c, ϕ_1, \dots, ϕ_p as well as the time series
- ▶ Note that $c, \phi_1, \dots, \phi_p, \theta_1, \dots, \theta_q$ are all unknown

Time-Series Forecasts: Estimating the ARIMA Unknowns

Assume we have
selected p and q
(and d)

- ▶ In the ARIMA model we can use MLE (default in the python library we use) to estimate $c, \phi_1, \dots, \phi_p, \theta_1, \dots, \theta_q$ which are all unknown
- ▶ For the MLE, we maximize the likelihood of obtaining the series y'_t for all t assuming it is Gaussian and so are the ϵ_t values
- ▶ Another approach is to minimize the sum of squares of ϵ_t
- ▶ Either case would result in a complicated non-linear optimization problem that cannot be written in closed-form //
- ▶ There are numerous software packages that produce the results

Time-Series Forecasts: Model for SARIMAX



- ▶ The key idea of seasonal-ARIMA is to add seasonal terms
- ▶ Extend the **ARIMA**(p, d, q) model to the **seasonal-ARIMA**(p, d, q)(P, D, Q) with m seasons
- ▶ Although we do not present the equation, it is worthwhile understanding that we add seasonality AR terms y_{t-m}, \dots, y_{t-pm} with coefficients Φ_1, \dots, Φ_P , as well as seasonality MA terms $\epsilon_{t-m}, \dots, \epsilon_{t-Qm}$ with coefficients $\Theta_1, \dots, \Theta_Q$
- ▶ Now, if we were to add exogenous variables, say x_1, x_2, \dots, x_r , then we essentially get the form

$$y'_t = c + b_1x_1 + b_2x_2 + \dots + b_rx_r + \phi_1y'_{t-1} + \dots$$

Selecting Model Parameters for SARIMAX

d, D by using pmdarima or simply plotting.

- ▶ Next we need to determine how to get $p, q, P,$ and Q
- ▶ They can be selected by searching through the space of values and find the one that minimizes an information criterion
- ▶ The common choice of information criteria are AIC (Akaike's Information Criterion), AICc (i.e. AIC corrected), and BIC (Bayesian Information Criterion), see (Hyndman and Athanasopoulos' FPP3 book)
- ▶ We do not present the expressions here but they are typically functions of likelihood of obtain the data (that was used in MLE), $T, p, q, P, Q,$ and whether or not c was used
- ▶ There is also an automated way (although many do not recommend) to do this using pmdarima

SARIMAX Forecasts: Python Code

- ▶ Import various python tools and libraries
- ▶ Select params viz. test data, sample size, etc.
- ▶ Get data and aggregate, if necessary
- ▶ Clean up data by imputing
- ▶ Merge with exogenous data (if we use them) to get required dataset
- ▶ Visualize the data for a few days
- ▶ Fit auto-arma model, or other methods for (p, d, q) and (P, D, Q)
- ▶ Use sliding window for predictions
- ▶ Test predictions using various metrics